

10)

a)

```
> library(ISLR)
```

```
> summary(Weekly)
```

Year	Lag1	Lag2	Lag3
Min. :1990	Min. :-18.1950	Min. :-18.1950	Min. :-18.1950
1st Qu.:1995	1st Qu.: -1.1540	1st Qu.: -1.1540	1st Qu.: -1.1580
Median :2000	Median : 0.2410	Median : 0.2410	Median : 0.2410
Mean :2000	Mean : 0.1506	Mean : 0.1511	Mean : 0.1472
3rd Qu.:2005	3rd Qu.: 1.4050	3rd Qu.: 1.4090	3rd Qu.: 1.4090
Max. :2010	Max. :12.0260	Max. :12.0260	Max. :12.0260

Lag4	Lag5	Volume	Today
Min. :-18.1950	Min. :-18.1950	Min. :0.08747	Min. :-18.1950
1st Qu.: -1.1580	1st Qu.: -1.1660	1st Qu.:0.33202	1st Qu.: -1.1540
Median : 0.2380	Median : 0.2340	Median :1.00268	Median : 0.2410
Mean : 0.1458	Mean : 0.1399	Mean :1.57462	Mean : 0.1499
3rd Qu.: 1.4090	3rd Qu.: 1.4050	3rd Qu.:2.05373	3rd Qu.: 1.4050
Max. :12.0260	Max. :12.0260	Max. :9.32821	Max. :12.0260

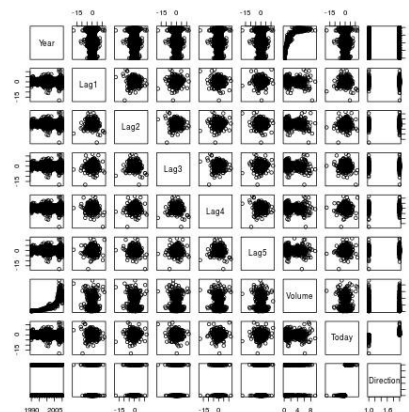
Direction
Down:484
Up :605

```
> jpeg("pairs-10a.jpg")
```

```
> pairs(Weekly)
```

```
> dev.off()
```

```
> cor(Weekly[, -9])
```



	Year	Lag1	Lag2	Lag3	Lag4
Year	1.00000000	-0.032289274	-0.03339001	-0.03000649	-0.031127923
Lag1	-0.03228927	1.000000000	-0.07485305	0.05863568	-0.071273876
Lag2	-0.03339001	-0.074853051	1.000000000	-0.07572091	0.058381535
Lag3	-0.03000649	0.058635682	-0.07572091	1.000000000	-0.075395865
Lag4	-0.03112792	-0.071273876	0.05838153	-0.07539587	1.000000000
Lag5	-0.03051910	-0.008183096	-0.07249948	0.06065717	-0.075675027
Volume	0.84194162	-0.064951313	-0.08551314	-0.06928771	-0.061074617
Today	-0.03245989	-0.075031842	0.05916672	-0.07124364	-0.007825873

	Lag5	Volume	Today
Year	-0.030519101	0.84194162	-0.032459894
Lag1	-0.008183096	-0.06495131	-0.075031842
Lag2	-0.072499482	-0.08551314	0.059166717
Lag3	0.060657175	-0.06928771	-0.071243639
Lag4	-0.075675027	-0.06107462	-0.007825873
Lag5	1.000000000	-0.05851741	0.011012698

```
Volume -0.058517414 1.000000000 -0.033077783
Today 0.011012698 -0.03307778 1.000000000
```

Year and Volume appear to be related. No other pattern is conceivable.

b)

```
> attach(Weekly)
```

The following objects are masked from Weekly (pos = 3):

Direction, Lag1, Lag2, Lag3, Lag4, Lag5, Today, Volume, Year

```
> glm.fit = glm(Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 + Volume, data = Weekly, family =
binomial)
```

```
> summary(glm.fit)
```

Call:

```
glm(formula = Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 +
    Volume, family = binomial, data = Weekly)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.6949	-1.2565	0.9913	1.0849	1.4579

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	0.26686	0.08593	3.106	0.0019 **
Lag1	-0.04127	0.02641	-1.563	0.1181
Lag2	0.05844	0.02686	2.175	0.0296 *
Lag3	-0.01606	0.02666	-0.602	0.5469
Lag4	-0.02779	0.02646	-1.050	0.2937
Lag5	-0.01447	0.02638	-0.549	0.5833
Volume	-0.02274	0.03690	-0.616	0.5377

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1496.2 on 1088 degrees of freedom
Residual deviance: 1486.4 on 1082 degrees of freedom
AIC: 1500.4

Number of Fisher Scoring iterations: 4

Lag2 appears to have some statistical significance with Pr = 0.0296

c)

```
> glm.probs = predict(glm.fit, type = "response")
```

```
> glm.pred = rep("Down", length(glm.probs))
```

```
> glm.pred[glm.probs > 0.5] = "Up"
> table(glm.pred, Direction)
      Direction
glm.pred Down Up
Down    54  48
Up     430 557
```

Percentage of current predictions: $(54+557)/(54+557+48+430) = 0.561$

Weeks the market goes up, the logistic regression is correct most of the time: $557/(557+48) = 0.921$

Weeks the market goes up, the logistic regression is wrong most of the time: $48/(557+48) = 0.0793$

d)

```
> train = (Year < 2009)
> Weekly.2009.2010 = Weekly[!train, ]
> glm.fit = glm(Direction ~ Lag2, data = Weekly, family = binomial, subset = train)
> glm.probs = predict(glm.fit, Weekly.2009.2010 , type = "response")
> glm.pred = rep("Down", length(glm.probs))
> glm.pred[glm.probs > 0.5] = "Up"
> Direction.2009.2010 = Direction[!train]
> table(glm.pred, Direction.2009.2010 )
      Direction.2009.2010
glm.pred Down Up
Down      9  5
Up      34 56
```

```
> mean(glm.pred == Direction.2009.2010 )
[1] 0.625
```

e)

```
> library(MASS)
> lda.fit = lda(Direction ~ Lag2, data = Weekly, subset = train)
> lda.pred = predict(lda.fit, Weekly.2009.2010 )
> table(lda.pred$class, Direction.2009.2010 )
      Direction.2009.2010
lda.pred Down Up
Down      9  5
Up      34 56
```

```
> mean(lda.pred$class == Direction.2009.2010 )
[1] 0.625
```

f)

```
> qda.fit = qda(Direction ~ Lag2, data = Weekly, subset = train)
> qda.class = predict(qda.fit, Weekly.2009.2010)$class
> table(qda.class, Direction.2009.2010)
      Direction.2009.2010
qda.class Down Up
```

```
Down 0 0
Up   43 61
```

```
> mean(qda.class == Direction.2009.2010)
[1] 0.5865385
```

The correctness is 0.58, with all predictions are up

```
g)
> library(class)
> train.X = as.matrix(Lag2[train])
> test.X = as.matrix(Lag2[!train])
> train.Direction = Direction[train]
> set.seed(1)
> knn.pred = knn(train.X, test.X, train.Direction, k = 1)
> table(knn.pred, Direction.2009.2010)
      Direction.2009.2010
knn.pred Down Up
      Down  21 30
       Up   22 31
```

```
> mean(knn.pred == Direction.2009.2010)
[1] 0.5
```

h)
Logistic regression and LDA provide similar test error rates

```
i)
> # Logistic regression with Lag2:Lag1
> glm.fit = glm(Direction ~ Lag2:Lag1, data = Weekly, family = binomial, subset = train)
> glm.probs = predict(glm.fit, Weekly.2009.2010, type = "response")
> glm.pred = rep("Down", length(glm.probs))
> glm.pred[glm.probs > 0.5] = "Up"
> Direction.2009.2010 = Direction[!train]
> table(glm.pred, Direction.2009.2010)
      Direction.2009.2010
glm.pred Down Up
      Down   1  1
       Up   42 60
```

```
> mean(glm.pred == Direction.2009.2010)
[1] 0.5865385
```

```
> # LDA with Lag2 interaction with Lag1
> lda.fit = lda(Direction ~ Lag2:Lag1, data = Weekly, subset = train)
> lda.pred = predict(lda.fit, Weekly.2009.2010)
> mean(lda.pred$class == Direction.2009.2010)
[1] 0.5769231
```

```

> # QDA with sqrt(abs(Lag2))
> qda.fit = qda(Direction ~ Lag2 + sqrt(abs(Lag2)), data = Weekly, subset = train)
> qda.class = predict(qda.fit, Weekly.2009.2010)$class
> table(qda.class, Direction.2009.2010)
      Direction.2009.2010
qda.class Down Up
      Down  12 13
      Up   31 48

> mean(qda.class == Direction.2009.2010)
[1] 0.5769231

> # KNN k = 10
> knn.pred = knn(train.X, test.X, train.Direction, k = 10)
> table(knn.pred, Direction.2009.2010)
      Direction.2009.2010
knn.pred Down Up
      Down  17 18
      Up   26 43

> mean(knn.pred == Direction.2009.2010)
[1] 0.5769231

> # KNN k = 100
> knn.pred = knn(train.X, test.X, train.Direction, k = 100)
> table(knn.pred, Direction.2009.2010)
      Direction.2009.2010
knn.pred Down Up
      Down   9 12
      Up   34 49

> mean(knn.pred == Direction.2009.2010)
[1] 0.5576923

```

Overall, the original LDA and logistic regression provide better performance in terms of test error rates

11)

a)

```

> library(ISLR)
> summary(Auto)
      mpg      cylinders  displacement  horsepower    weight
Min.   :9.00  Min.   :3.000  Min.   :68.0  Min.   :46.0  Min.   :1613
1st Qu.:17.00 1st Qu.:4.000 1st Qu.:105.0 1st Qu.: 75.0 1st Qu.:2225
Median :22.75 Median :4.000 Median :151.0 Median : 93.5 Median :2804
Mean   :23.45 Mean   :5.472 Mean  :194.4 Mean  :104.5 Mean  :2978
3rd Qu.:29.00 3rd Qu.:8.000 3rd Qu.:275.8 3rd Qu.:126.0 3rd Qu.:3615
Max.   :46.60 Max.   :8.000 Max.   :455.0 Max.   :230.0 Max.   :5140

```

acceleration	year	origin	name	
Min. :8.00	Min. :70.00	Min. :1.000	amc matador	: 5
1st Qu.:13.78	1st Qu.:73.00	1st Qu.:1.000	ford pinto	: 5
Median :15.50	Median :76.00	Median :1.000	toyota corolla	: 5
Mean :15.54	Mean :75.98	Mean :1.577	amc gremlin	: 4
3rd Qu.:17.02	3rd Qu.:79.00	3rd Qu.:2.000	amc hornet	: 4
Max. :24.80	Max. :82.00	Max. :3.000	chevrolet chevette:	4
	(Other)			:365

```
> mpg01 = rep(0, length(mpg))
> mpg01[mpg > median(mpg)] = 1
> Auto = data.frame(Auto, mpg01)
```

```
> head(Auto)
  mpg cylinders displacement horsepower weight acceleration year origin
1  18         8        307        130  3504         12.0  70     1
2  15         8        350        165  3693         11.5  70     1
3  18         8        318        150  3436         11.0  70     1
4  16         8        304        150  3433         12.0  70     1
5  17         8        302        140  3449         10.5  70     1
6  15         8        429        198  4341         10.0  70     1

      name mpg01 mpg01.1
1 chevrolet chevelle malibu    0    0
2   buick skylark 320    0    0
3  plymouth satellite    0    0
4    amc rebel sst    0    0
5    ford torino    0    0
6  ford galaxie 500    0    0
```

b)

```
> cor(Auto[, -9])

      mpg cylinders displacement horsepower weight
mpg      1.0000000 -0.7776175 -0.8051269 -0.7784268 -0.8322442
cylinders -0.7776175  1.0000000  0.9508233  0.8429834  0.8975273
displacement -0.8051269  0.9508233  1.0000000  0.8972570  0.9329944
horsepower -0.7784268  0.8429834  0.8972570  1.0000000  0.8645377
weight      -0.8322442  0.8975273  0.9329944  0.8645377  1.0000000
acceleration 0.4233285 -0.5046834 -0.5438005 -0.6891955 -0.4168392
year         0.5805410 -0.3456474 -0.3698552 -0.4163615 -0.3091199
origin       0.5652088 -0.5689316 -0.6145351 -0.4551715 -0.5850054
mpg01        0.8369392 -0.7591939 -0.7534766 -0.6670526 -0.7577566
mpg01.1       0.8369392 -0.7591939 -0.7534766 -0.6670526 -0.7577566

      acceleration year origin mpg01 mpg01.1
mpg      0.4233285 0.5805410 0.5652088 0.8369392 0.8369392
cylinders -0.5046834 -0.3456474 -0.5689316 -0.7591939 -0.7591939
displacement -0.5438005 -0.3698552 -0.6145351 -0.7534766 -0.7534766
horsepower -0.6891955 -0.4163615 -0.4551715 -0.6670526 -0.6670526
weight      -0.4168392 -0.3091199 -0.5850054 -0.7577566 -0.7577566
```

```

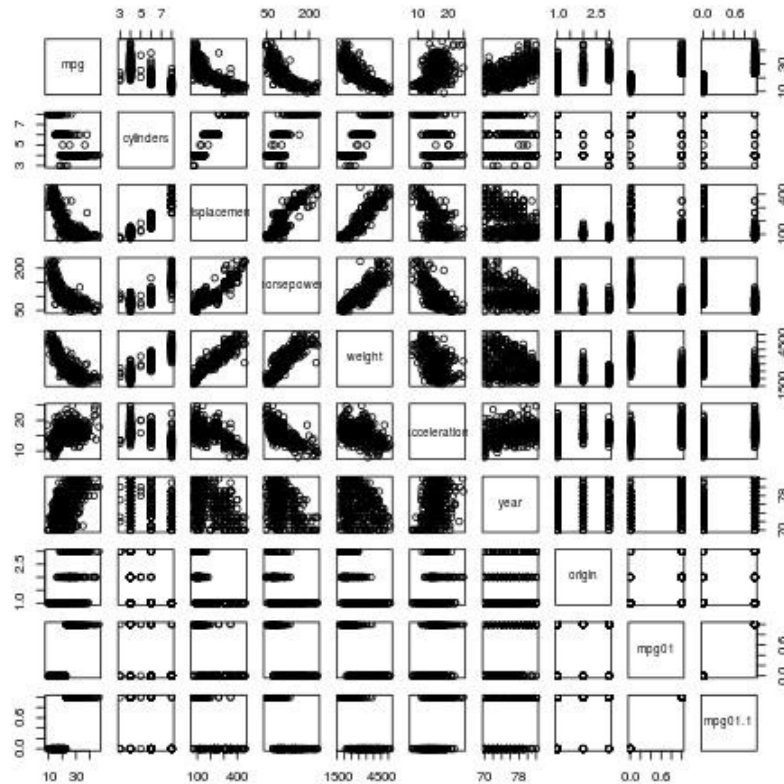
acceleration  1.0000000 0.2903161 0.2127458 0.3468215 0.3468215
year          0.2903161 1.0000000 0.1815277 0.4299042 0.4299042
origin        0.2127458 0.1815277 1.0000000 0.5136984 0.5136984
mpg01         0.3468215 0.4299042 0.5136984 1.0000000 1.0000000
mpg01.1       0.3468215 0.4299042 0.5136984 1.0000000 1.0000000

```

```

> jpeg("pairs-11b")
> pairs(Auto[, -9])
> dev.off()

```



Having a negative correlation with cylinders, weight, displacement, horsepower and mpg.

```

c)
> train = (year%%2 == 0) # if the year is even, data is training
> test = !train
> Auto.train = Auto[train, ]
> Auto.test = Auto[test, ]
> mpg01.test = mpg01[test]

```

```

d)
> # LDA
> library(MASS)
> lda.fit = lda(mpg01 ~ cylinders + weight + displacement + horsepower, data = Auto, subset = train)
> lda.pred = predict(lda.fit, Auto.test)
> mean(lda.pred$class != mpg01.test)
[1] 0.1263736

```

Test error rate: 0.1263736

e)

```
> # QDA
> qda.fit = qda(mpg01 ~ cylinders + weight + displacement + horsepower, data = Auto, subset = train)
> qda.pred = predict(qda.fit, Auto.test)
> mean(qda.pred$class != mpg01.test)
[1] 0.1318681
```

Test error rate: 0.1318681

f)

```
> # Logistic regression
> glm.fit = glm(mpg01 ~ cylinders + weight + displacement + horsepower, data = Auto, family =
binomial, subset = train)
> glm.probs = predict(glm.fit, Auto.test, type = "response")
> glm.pred = rep(0, length(glm.probs))
> glm.pred[glm.probs > 0.5] = 1
> mean(glm.pred != mpg01.test)
[1] 0.1208791
```

Test error rate: 0.1208791

g)

```
> library(class)
> train.X = cbind(cylinders, weight, displacement, horsepower)[train, ]
> test.X = cbind(cylinders, weight, displacement, horsepower)[test, ]
> train.mpg01 = mpg01[train]
> set.seed(1)
> # KNN(k=1)
> knn.pred = knn(train.X, test.X, train.mpg01, k = 1)
> mean(knn.pred != mpg01.test)
[1] 0.1538462
```

```
> # KNN(k=10)
> knn.pred = knn(train.X, test.X, train.mpg01, k = 10)
> mean(knn.pred != mpg01.test)
[1] 0.1648352
```

```
> # KNN(k=100)
> knn.pred = knn(train.X, test.X, train.mpg01, k = 100)
> mean(knn.pred != mpg01.test)
[1] 0.1428571
```

Test error rate (k=1): 0.1538462

Test error rate (k=10): 0.1648352

Test error rate (k=100): 0.1428571

with k=100, KNN produces the most significant performance.