


2)

g)

```
> jpeg("plot_2g.jpg")
> pr = function(n) return(1 - (1 - 1/n)^n)
> x = 1:1e+05
> plot(x, pr(x))
> dev.off()
```

the plot reaches asymptote of 0.63



h)

```
> store=rep(NA, 10000)
> for(i in 1:10000) {
+ store[i]=sum(sample(1:100,
rep=TRUE)==4) > 0
+ }
> mean(store)
[1] 0.6331
```

This confirms the results found in the previous part

5)

a)

```
> library(ISLR)
> summary(Default)
default student balance income
No :9667 No :7056 Min. : 0.0 Min. : 772
Yes: 333 Yes:2944 1st Qu.: 481.7 1st Qu.:21340
Median : 823.6 Median :34553
Mean : 835.4 Mean :33517
3rd Qu.:1166.3 3rd Qu.:43808
Max. :2654.3 Max. :73554

> attach(Default)
> set.seed(1)
> glm.fit = glm(default ~ income + balance, data = Default, family = binomial)
> glm.fit
```

Call: glm(formula = default ~ income + balance, family = binomial,
data = Default)

Coefficients:

```
(Intercept) income balance
-1.154e+01 2.081e-05 5.647e-03
```

Degrees of Freedom: 9999 Total (i.e. Null); 9997 Residual

Null Deviance: 2921

Residual Deviance: 1579 AIC: 1585

```

b)
# i.
> train = sample(dim(Default)[1], dim(Default)[1]/2)
# ii.
> glm.fit = glm(default ~ income + balance, data = Default, family = binomial, subset = train)
# iii.
> glm.pred = rep("No", dim(Default)[1]/2)
> glm.probs = predict(glm.fit, Default[-train, ], type = "response")
> glm.pred[glm.probs > 0.5] = "Yes"
# iv.
> mean(glm.pred != Default[-train, ]$default)
[1] 0.0286

```

Test error rate from validation set approach: 0.0286

```

c)
#1st
[1] 0.0236

```

```

#2nd
[1] 0.0252

```

```

#3rd
[1] 0.0268

```

On average, test error rate: 0.02

```

d)
> train = sample(dim(Default)[1], dim(Default)[1]/2)
> glm.fit = glm(default ~ income + balance + student, data = Default, family = binomial,
+ subset = train)
> glm.pred = rep("No", dim(Default)[1]/2)
> glm.probs = predict(glm.fit, Default[-train, ], type = "response")
> glm.pred[glm.probs > 0.5] = "Yes"
> mean(glm.pred != Default[-train, ]$default)
[1] 0.0264

```

Test error rate: 0.0264 with dummy variable ‘student’, it doesn’t appear to reduce the test error rate with adding dummy variable ‘student’

```

7)
a)
> library(ISLR)
> summary(Weekly)
      Year      Lag1      Lag2      Lag3
Min. :1990 Min. :-18.1950 Min. :-18.1950 Min. :-18.1950
1st Qu.:1995 1st Qu.: -1.1540 1st Qu.: -1.1540 1st Qu.: -1.1580

```

```

Median :2000  Median : 0.2410  Median : 0.2410  Median : 0.2410
Mean   :2000  Mean   : 0.1506  Mean   : 0.1511  Mean   : 0.1472
3rd Qu.:2005  3rd Qu.: 1.4050  3rd Qu.: 1.4090  3rd Qu.: 1.4090
Max.   :2010  Max.   :12.0260  Max.   :12.0260  Max.   :12.0260
  Lag4      Lag5      Volume      Today
Min.   :-18.1950  Min.   :-18.1950  Min.   :0.08747  Min.   :-18.1950
1st Qu.: -1.1580  1st Qu.: -1.1660  1st Qu.:0.33202  1st Qu.: -1.1540
Median : 0.2380  Median : 0.2340  Median :1.00268  Median : 0.2410
Mean   : 0.1458  Mean   : 0.1399  Mean   :1.57462  Mean   : 0.1499
3rd Qu.: 1.4090  3rd Qu.: 1.4050  3rd Qu.:2.05373  3rd Qu.: 1.4050
Max.   :12.0260  Max.   :12.0260  Max.   :9.32821  Max.   :12.0260
Direction
Down:484
Up :605

```

```

> set.seed(1)
> attach(Weekly)
> glm.fit = glm(Direction ~ Lag1 + Lag2, data = Weekly, family = binomial)
> summary(glm.fit)

```

Call:

```
glm(formula = Direction ~ Lag1 + Lag2, family = binomial, data = Weekly)
```

Deviance Residuals:

```

    Min      1Q  Median      3Q     Max
-1.623 -1.261  1.001  1.083  1.506

```

Coefficients:

```

            Estimate Std. Error z value Pr(>|z|)
(Intercept) 0.22122   0.06147   3.599 0.000319 ***
Lag1        -0.03872   0.02622  -1.477 0.139672
Lag2         0.06025   0.02655   2.270 0.023232 *
---

```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

```

Null deviance: 1496.2 on 1088 degrees of freedom
Residual deviance: 1488.2 on 1086 degrees of freedom
AIC: 1494.2

```

Number of Fisher Scoring iterations: 4

b)

```

> glm.fit = glm(Direction ~ Lag1 + Lag2, data = Weekly[-1, ], family = binomial)
> summary(glm.fit)

```

Call:

```
glm(formula = Direction ~ Lag1 + Lag2, family = binomial, data = Weekly[-1, ])
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.6258	-1.2617	0.9999	1.0819	1.5071

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	0.22324	0.06150	3.630	0.000283 ***
Lag1	-0.03843	0.02622	-1.466	0.142683
Lag2	0.06085	0.02656	2.291	0.021971 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1494.6 on 1087 degrees of freedom
Residual deviance: 1486.5 on 1085 degrees of freedom
AIC: 1492.5

Number of Fisher Scoring iterations: 4

c)

```
> predict.glm(glm.fit, Weekly[1, ], type = "response")
1
0.5713923
> Weekly[1, ]
  Year Lag1 Lag2 Lag3 Lag4 Lag5 Volume Today Direction
1 1990 0.816 1.572 -3.936 -0.229 -3.484 0.154976 -0.27   Down
```

So the prediction is not correct, since prediction is UP, true direction was DOWN.

d)

```
> count = rep(0, dim(Weekly)[1])
> for (i in 1:(dim(Weekly)[1])) {
+   glm.fit = glm(Direction ~ Lag1 + Lag2, data = Weekly[-i, ], family = binomial)
+   is_up = predict.glm(glm.fit, Weekly[i, ], type = "response") > 0.5
+   is_true_up = Weekly[i, ]$Direction == "Up"
+   if (is_up != is_true_up)
+     count[i] = 1
+ }
> sum(count)
[1] 490
```

Total errors: 490

e)

```
> mean(count)
[1] 0.4499541
LOOCV estimates a test error rate of 0.45
```

9)

a)

```
> library(MASS)
```

```
> summary(Boston)
```

crim	zn	indus	chas
Min. : 0.00632	Min. : 0.00	Min. : 0.46	Min. : 0.00000
1st Qu.: 0.08204	1st Qu.: 0.00	1st Qu.: 5.19	1st Qu.: 0.00000
Median : 0.25651	Median : 0.00	Median : 9.69	Median : 0.00000
Mean : 3.61352	Mean : 11.36	Mean : 11.14	Mean : 0.06917
3rd Qu.: 3.67708	3rd Qu.: 12.50	3rd Qu.: 18.10	3rd Qu.: 0.00000
Max. : 88.97620	Max. : 100.00	Max. : 27.74	Max. : 1.00000

nox	rm	age	dis
Min. : 0.3850	Min. : 3.561	Min. : 2.90	Min. : 1.130
1st Qu.: 0.4490	1st Qu.: 5.886	1st Qu.: 45.02	1st Qu.: 2.100
Median : 0.5380	Median : 6.208	Median : 77.50	Median : 3.207
Mean : 0.5547	Mean : 6.285	Mean : 68.57	Mean : 3.795
3rd Qu.: 0.6240	3rd Qu.: 6.623	3rd Qu.: 94.08	3rd Qu.: 5.188
Max. : 0.8710	Max. : 8.780	Max. : 100.00	Max. : 12.127

rad	tax	ptratio	black
Min. : 1.000	Min. : 187.0	Min. : 12.60	Min. : 0.32
1st Qu.: 4.000	1st Qu.: 279.0	1st Qu.: 17.40	1st Qu.: 375.38
Median : 5.000	Median : 330.0	Median : 19.05	Median : 391.44
Mean : 9.549	Mean : 408.2	Mean : 18.46	Mean : 356.67
3rd Qu.: 24.000	3rd Qu.: 666.0	3rd Qu.: 20.20	3rd Qu.: 396.23
Max. : 24.000	Max. : 711.0	Max. : 22.00	Max. : 396.90

lstat	medv
Min. : 1.73	Min. : 5.00
1st Qu.: 6.95	1st Qu.: 17.02
Median : 11.36	Median : 21.20
Mean : 12.65	Mean : 22.53
3rd Qu.: 16.95	3rd Qu.: 25.00
Max. : 37.97	Max. : 50.00

```
> mean(medv)
```

```
[1] 22.53281
```

b)

```
> sd(medv)/sqrt(length(medv))
```

```
[1] 0.4088611
```

c)

```
> boot(medv, function(data, index) return(mean(data[index])), 1000)
```

ORDINARY NONPARAMETRIC BOOTSTRAP

Call:

```
boot(data = medv, statistic = function(data, index) return(mean(data[index])),  
      R = 1000)
```

Bootstrap Statistics :

	original	bias	std. error
t1*	22.53281	0.0027583	0.4120131

Standard error: 0.4120131 similar to part b) 0.4088611

d)

```
> t.test(medv)
```

One Sample t-test

data: medv

t = 55.111, df = 505, p-value < 2.2e-16

alternative hypothesis: true mean is not equal to 0

95 percent confidence interval:

21.72953 23.33608

sample estimates:

mean of x

22.53281

```
> c(bstrap$t0 - 2 * 0.4119, bstrap$t0 + 2 * 0.4119)
```

```
[1] 21.70901 23.35661
```

Bootstrap estimates (21.70901, 23.35661) compared to (21.72953, 23.33608) t-test

e)

```
> median(medv)
```

```
[1] 21.2
```

f)

```
> boot.fn = function(data, index) return(median(data[index]))
```

```
> boot(medv, boot.fn, 1000)
```

ORDINARY NONPARAMETRIC BOOTSTRAP

Call:

```
boot(data = medv, statistic = boot.fn, R = 1000)
```

Bootstrap Statistics :

	original	bias	std. error
t1*	21.2	-0.0025	0.374358

Bootstrap estimates median of 21.2 with standard error: 0.374358

```
g)
> quantile(medv, c(0.1))
10%
12.75
```

```
h)
> boot.fn = function(data, index) return(quantile(data[index], c(0.1)))
> boot(medv, boot.fn, 1000)
```

ORDINARY NONPARAMETRIC BOOTSTRAP

Call:
boot(data = medv, statistic = boot.fn, R = 1000)

Bootstrap Statistics :

	original	bias	std. error
t1*	12.75	0.0261	0.4912231

Bootstrap estimates 10th-quantile of 12.75 with standard error 0.4912231