

```

9)
a)
> set.seed(1)
> library(MASS)
> attach(Boston)
> lm.fit = lm(nox ~ poly(dis, 3), data = Boston)
> summary(lm.fit)

```

Call:

```
lm(formula = nox ~ poly(dis, 3), data = Boston)
```

Residuals:

| Min | 1Q | Median | 3Q | Max |
|-----------|-----------|-----------|----------|----------|
| -0.121130 | -0.040619 | -0.009738 | 0.023385 | 0.194904 |

Coefficients:

| | Estimate | Std. Error | t value | Pr(> t) |
|---------------|-----------|------------|---------|--------------|
| (Intercept) | 0.554695 | 0.002759 | 201.021 | < 2e-16 *** |
| poly(dis, 3)1 | -2.003096 | 0.062071 | -32.271 | < 2e-16 *** |
| poly(dis, 3)2 | 0.856330 | 0.062071 | 13.796 | < 2e-16 *** |
| poly(dis, 3)3 | -0.318049 | 0.062071 | -5.124 | 4.27e-07 *** |

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.06207 on 502 degrees of freedom

Multiple R-squared: 0.7148, Adjusted R-squared: 0.7131


F-statistic: 419.3 on 3 and 502 DF, p-value: < 2.2e-16

```

> dislim = range(dis)
> dis.grid = seq(from = dislim[1], to = dislim[2], by = 0.1)
> lm.pred = predict(lm.fit, list(dis = dis.grid))
> jpeg("9a.jpg")
> plot(nox ~ dis, data = Boston, col = "darkgrey")
> lines(dis.grid, lm.pred, col = "red", lwd = 2)
> dev.off()

```

Summary shows that all polynomial terms are significant while predicting nox using dis. Plot shows a smooth curve fitting the data fairly well.



b)

We plot polynomials of degrees 1 to 10 and save train RSS.

```

> all.rss = rep(NA, 10)
> for (i in 1:10) {
+   lm.fit = lm(nox ~ poly(dis, i), data = Boston)
+   all.rss[i] = sum(lm.fit$residuals^2)
+ }
> all.rss
[1] 2.768563 2.035262 1.934107 1.932981 1.915290 1.878257 1.849484 1.835630

```


[9] 1.833331 1.832171

As expected, train RSS monotonically decreases with degree of polynomial.

c)

We use a 10-fold cross validation to pick the best polynomial degree.

```
> jpeg("9b.jpg")
> plot(1:10, all.deltas, xlab = "Degree", ylab = "CV
error", type = "l", pch = 20, lwd = 2)
> dev.off()
```



A 10-fold CV shows that the CV error reduces as we increase degree from 1 to 3, stay almost constant till degree 5, and the starts increasing for higher degrees. We pick 4 as the best polynomial degree.

d)

We see that dis has limits of about 1 and 13 respectively. We split this range in roughly equal 4 intervals and establish knots at [4,7,11]. Note: bs function in R expects either df or knots argument. If both are specified, knots are ignored.

```
> library(splines)
> sp.fit = lm(nox ~ bs(dis, df = 4, knots = c(4, 7, 11)), data = Boston)
> summary(sp.fit)
```

Call:

```
lm(formula = nox ~ bs(dis, df = 4, knots = c(4, 7, 11)), data = Boston)
```

Residuals:

| Min | 1Q | Median | 3Q | Max |
|-----------|-----------|-----------|----------|----------|
| -0.124567 | -0.040355 | -0.008702 | 0.024740 | 0.192920 |

Coefficients:

| | Estimate | Std. Error | t value | Pr(> t) |
|---------------------------------------|----------|------------|---------|--------------|
| (Intercept) | 0.73926 | 0.01331 | 55.537 | < 2e-16 *** |
| bs(dis, df = 4, knots = c(4, 7, 11))1 | -0.08861 | 0.02504 | -3.539 | 0.00044 *** |
| bs(dis, df = 4, knots = c(4, 7, 11))2 | -0.31341 | 0.01680 | -18.658 | < 2e-16 *** |
| bs(dis, df = 4, knots = c(4, 7, 11))3 | -0.26618 | 0.03147 | -8.459 | 3.00e-16 *** |
| bs(dis, df = 4, knots = c(4, 7, 11))4 | -0.39802 | 0.04647 | -8.565 | < 2e-16 *** |
| bs(dis, df = 4, knots = c(4, 7, 11))5 | -0.25681 | 0.09001 | -2.853 | 0.00451 ** |
| bs(dis, df = 4, knots = c(4, 7, 11))6 | -0.32926 | 0.06327 | -5.204 | 2.85e-07 *** |

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.06185 on 499 degrees of freedom

Multiple R-squared: 0.7185, Adjusted R-squared: 0.7151


F-statistic: 212.3 on 6 and 499 DF, p-value: < 2.2e-16

```

> sp.pred = predict(sp.fit, list(dis = dis.grid))
> jpeg("9c.jpg")
> plot(nox ~ dis, data = Boston, col = "darkgrey")
> lines(dis.grid, sp.pred, col = "red", lwd = 2)
> dev.off()

```

The summary shows that all terms in spline fit are significant. Plot shows that the spline fits data well except at the extreme values of dis, (especially $dis > 10$).



e)

We fit regression splines with dfs between 3 and 16.

```

> all.cv = rep(NA, 16)
> for (i in 3:16) {
+   lm.fit = lm(nox ~ bs(dis, df = i), data = Boston)
+   all.cv[i] = sum(lm.fit$residuals^2)
+ }
> all.cv[-c(1, 2)]
[1] 1.934107 1.922775 1.840173 1.833966 1.829884 1.816995 1.825653 1.792535
[9] 1.796992 1.788999 1.782350 1.781838 1.782798 1.783546

```

Train RSS monotonically decreases till $df=14$ and then slightly increases for $df=15$ and $df=16$.

f)

Finally, we use a 10-fold cross validation to find best df.
We try all integer values of df between 3 and 16.

```

> all.cv = rep(NA, 16)
> for (i in 3:16) {
+   lm.fit = glm(nox ~ bs(dis, df = i), data = Boston)
+   all.cv[i] = cv.glm(Boston, lm.fit, K = 10)$delta[2]
+ }


```

```

> jpeg("9f.jpg")
> plot(3:16, all.cv[-c(1, 2)], lwd = 2, type = "l", xlab =
"df", ylab = "CV error")
> dev.off()

```

CV error is more jumpy in this case, but attains minimum at $df=10$. We pick 1010 as the optimal degrees of freedom.



11)

a)

We create variables according to the equation $Y = -2.1 + 1.3X_1 + 0.54X_2$.

```
> set.seed(1)
> X1 = rnorm(100)
> X2 = rnorm(100)
> eps = rnorm(100, sd = 0.1)
> Y = -2.1 + 1.3 * X1 + 0.54 * X2 + eps
```

b)

Create a list of 1000 β^0 , β^1 and β^2 . Initialize first of the β^1 to 10.

```
> beta0 = rep(NA, 1000)
> beta1 = rep(NA, 1000)
> beta2 = rep(NA, 1000)
> beta1[1] = 10
```


c, d, e)

Accumulate results of 1000 iterations in the beta arrays.

```
> for (i in 1:1000) {
+   a = Y - beta1[i] * X1
+   beta2[i] = lm(a ~ X2)$coef[2]
+   a = Y - beta2[i] * X2
+   lm.fit = lm(a ~ X1)

+   if (i < 1000) {+     beta1[i + 1] =
lm.fit$coef[2]
+   }
+   beta0[i] = lm.fit$coef[1]
+ }

> jpeg("9cde.jpg")
> plot(1:1000, beta0, type = "l", xlab = "iteration",
ylab = "betas", ylim = c(-2.2,
+   1.6), col = "green")
> lines(1:1000, beta1, col = "red")
> lines(1:1000, beta2, col = "blue")
> legend("center", c("beta0", "beta1", "beta2"), lty
= 1, col = c("green", "red",
+   "blue"))
> dev.off()
```



The coefficients quickly attain their least square values.

f)

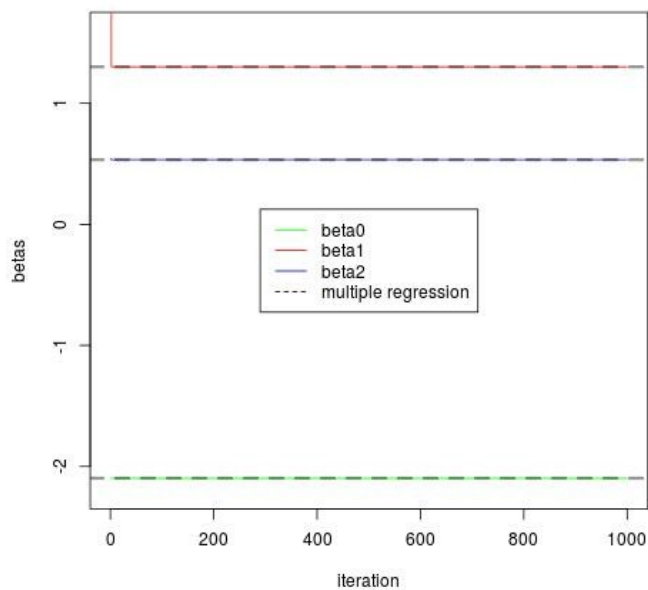
```
> lm.fit = lm(Y ~ X1 + X2)
```

```

> jpeg("9f.jpg")
> plot(1:1000, beta0, type = "l", xlab = "iteration", ylab = "betas", ylim = c(-2.2,
+ 1.6), col = "green")
> lines(1:1000, beta1, col = "red")
> lines(1:1000, beta2, col = "blue")
> abline(h = lm.fit$coef[1], lty = "dashed", lwd = 3, col = rgb(0, 0, 0, alpha = 0.4))
> abline(h = lm.fit$coef[2], lty = "dashed", lwd = 3, col = rgb(0, 0, 0, alpha = 0.4))
> abline(h = lm.fit$coef[3], lty = "dashed", lwd = 3, col = rgb(0, 0, 0, alpha = 0.4))
> legend("center", c("beta0", "beta1", "beta2", "multiple regression"), lty = c(1,
+ 1, 1, 2), col = c("green", "red", "blue", "black"))
> dev.off()

```

Dotted lines show that the estimated multiple regression coefficients match exactly with the coefficients obtained using backfitting.



g)
When the relationship between Y and X's is linear, one iteration is sufficient to attain a good approximation of true regression coefficients.