


5)
Majority approach

```
> p = c(0.1, 0.15, 0.2, 0.2, 0.55, 0.6, 0.6, 0.65, 0.7, 0.75)>
sum(p >= 0.5) > sum(p < 0.5)
[1] TRUE
```

The number of red predictions is greater than the number of green predictions based on a 50% threshold, thus RED.

Average approach
> mean(p)
[1] 0.45

The average of the probabilities is less than the 50% threshold, thus GREEN.



```
8)
a)
library(ISLR)
attach(Carseats)
set.seed(1)

train = sample(dim(Carseats)[1], dim(Carseats)[1]/2)
Carseats.train = Carseats[train, ]
Carseats.test = Carseats[-train, ]
```


```
b)
library(tree)
tree.carseats = tree(Sales ~ ., data = Carseats.train)
summary(tree.carseats)
##
## Regression tree:
## tree(formula = Sales ~ ., data = Carseats.train)
## Variables actually used in tree construction:
## [1] "ShelveLoc" "Price" "Age" "Advertising" "Income"
## [6] "CompPrice"
```

```
## Number of terminal nodes: 18 ## Residual mean
deviance: 2.36 = 429 / 182
## Distribution of residuals:
## Min. 1st Qu. Median Mean 3rd Qu. Max.
## -4.260 -1.040 0.102 0.000 0.930 3.910
plot(tree.carseats)
text(tree.carseats, pretty = 0)
pred.carseats = predict(tree.carseats, Carseats.test)
mean((Carseats.test$Sales - pred.carseats)^2)
## [1] 4.149
The test MSE is about 4.15
```

```
c)
cv.carseats = cv.tree(tree.carseats, FUN = prune.tree)
par(mfrow = c(1, 2))
plot(cv.carseats$size, cv.carseats$dev, type = "b")
plot(cv.carseats$k, cv.carseats$dev, type = "b")
```

```
# Best size = 9
pruned.carseats = prune.tree(tree.carseats, best = 9)
par(mfrow = c(1, 1))
plot(pruned.carseats)
text(pruned.carseats, pretty = 0)
```

```
pred.pruned = predict(pruned.carseats, Carseats.test)
mean((Carseats.test$Sales - pred.pruned)^2)
```



```
## [1] 4.993
```

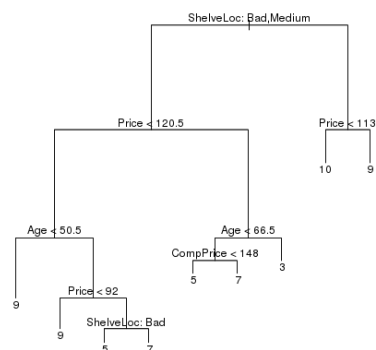
Pruning the tree in this case increases the test MSE to 4.99

d)

```
library(randomForest)
bag.carseats = randomForest(Sales ~ ., data = Carseats.train,
mtry = 10, ntree = 500, importance = T)
bag.pred = predict(bag.carseats, Carseats.test)
mean((Carseats.test$Sales - bag.pred)^2)
```

```
## [1] 2.586
```

```
importance(bag.carseats)
##          %IncMSE IncNodePurity
## CompPrice 13.8790   131.095
## Income    5.6042    77.033
## Advertising 14.1720  129.218
## Population 0.6071   65.196
## Price     53.6119  506.604
## ShelfLoc  44.0311  323.189
## Age       20.1751  189.269
## Education 1.4244   41.811
## Urban     -1.9640   8.124
## US        5.6989   14.307
```



Bagging improves the test MSE to 2.58. We also see that PricePrice, ShelfLocc and Age are three most important predictors of Sale.

e)

```
rf.carseats = randomForest(Sales ~ ., data = Carseats.train, mtry = 5, ntree = 500,
importance = T)
rf.pred = predict(rf.carseats, Carseats.test)
mean((Carseats.test$Sales - rf.pred)^2)
```

```
## [1] 2.87
```

```
importance(rf.carseats)
##          %IncMSE IncNodePurity
## CompPrice 11.2746   126.64
## Income    4.4397   101.63
## Advertising 12.9346  137.96
## Population 0.2725   78.78
## Price     49.2418  449.52
## ShelfLoc  38.8406  283.46
## Age       19.1329  195.14
## Education 1.9818   54.26
## Urban     -2.2083   11.35
## US        6.6487   26.71
```

In this case, random forest worsens the MSE on test set to 2.87. Changing mm varies test MSE between 2.6 to 3. We again see that Price, ShelfLoc and Age are three most important predictors of Sale.

9)

a)

```
library(ISLR)
attach(OJ)
set.seed(1013)
```

```
train = sample(dim(OJ)[1], 800)
OJ.train = OJ[train, ]
OJ.test = OJ[-train, ]
```

b)

```
library(tree)
```

```

oj.tree = tree(Purchase ~ ., data = OJ.train)
summary(oj.tree)
##
## Classification tree:
## tree(formula = Purchase ~ ., data = OJ.train)
## Variables actually used in tree construction:
## [1] "LoyalCH" "PriceDiff"
## Number of terminal nodes: 7
## Residual mean deviance: 0.752 = 596 / 793
## Misclassification error rate: 0.155 = 124 / 800

```

The tree only uses two variables: LoyalCH and PriceDiff. It has 7 terminal nodes. Training error rate (misclassification error) for the tree is 0.155

```

c)
oj.tree
## node), split, n, deviance, yval, (yprob)
##      * denotes terminal node
##
## 1) root 800 1000 CH ( 0.60 0.40 )
##    2) LoyalCH < 0.5036 359 400 MM ( 0.28 0.72 )
##      4) LoyalCH < 0.276142 170 100 MM ( 0.11 0.89 ) *
##      5) LoyalCH > 0.276142 189 300 MM ( 0.42 0.58 )
##        10) PriceDiff < 0.05 79 80 MM ( 0.19 0.81 ) *
##        11) PriceDiff > 0.05 110 100 CH ( 0.59 0.41 ) *
##    3) LoyalCH > 0.5036 441 300 CH ( 0.87 0.13 )
##      6) LoyalCH < 0.764572 186 200 CH ( 0.75 0.25 )
##        12) PriceDiff < -0.165 29 30 MM ( 0.28 0.72 ) *
##        13) PriceDiff > -0.165 157 100 CH ( 0.83 0.17 )
##          26) PriceDiff < 0.265 82 100 CH ( 0.73 0.27 ) *
##          27) PriceDiff > 0.265 75 30 CH ( 0.95 0.05 ) *
##    7) LoyalCH > 0.764572 255 90 CH ( 0.96 0.04 ) *

```

Let's pick terminal node labeled "10)". The splitting variable at this node is PriceDiff. The splitting value of this node is 0.05. There are 79 points in the subtree below this node. The deviance for all points contained in region below this node is 80. A * in the line denotes that this is in fact a terminal node. The prediction at this node is Sales = MM. About 19% points in this node have CH as value of Sales. Remaining 81% points have MM as value of Sales.

```

d)
plot(oj.tree)
text(oj.tree, pretty = 0)

```

LoyalCH is the most important variable of the tree, in fact top 3 nodes contain LoyalCH. If LoyalCH<0.27, the tree predicts MM. If LoyalCH>0.76, the tree predicts CH. For intermediate values of LoyalCH, the decision also depends on the value of PriceDiff.

```

e)
oj.pred = predict(oj.tree, OJ.test, type = "class")
table(OJ.test$Purchase, oj.pred)

```

```

##    oj.pred
##    CH MM
## CH 152 19
## MM 32 67

```

```

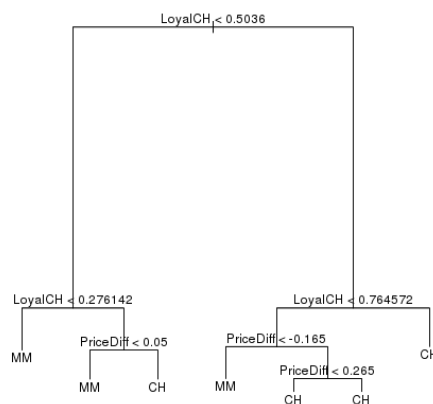
f)
cv.oj = cv.tree(oj.tree, FUN = prune.tree)

```

```

g)
plot(cv.oj$size, cv.oj$dev, type = "b", xlab = "Tree Size", ylab = "Deviance")

```




h) Size of 6 gives lowest cross-validation error.

i)
oj.pruned = prune.tree(oj.tree, best = 6)

j)
summary(oj.pruned)

```
##
## Classification tree:
## snip.tree(tree = oj.tree, nodes = 13L)
## Variables actually used in tree construction:
## [1] "LoyalCH" "PriceDiff"
## Number of terminal nodes: 6
## Residual mean deviance: 0.769 = 610 / 794
## Misclassification error rate: 0.155 = 124 / 800
```



Misclassification error of pruned tree is exactly same as that of original tree — 0.155.

k)
pred.unpruned = predict(oj.tree, OJ.test, type = "class")
misclass.unpruned = sum(OJ.test\$Purchase != pred.unpruned)
misclass.unpruned/length(pred.unpruned)


```
## [1] 0.1889
pred.pruned = predict(oj.pruned, OJ.test, type = "class")
misclass.pruned = sum(OJ.test$Purchase != pred.pruned)
misclass.pruned/length(pred.pruned)
```

```
## [1] 0.1889
Pruned and unpruned trees have same test error rate of 0.189.
```

l1)
a)
library(ISLR)
train = 1:1000
Caravan\$Purchase = ifelse(Caravan\$Purchase == "Yes", 1, 0)
Caravan.train = Caravan[train,]

Caravan.test = Caravan[-train,]

b)
library(gbm)
Loading required package: survival
Loading required package: splines
Loading required package: lattice
Loading required package: parallel
Loaded gbm 2.1
set.seed(342)
boost.caravan = gbm(Purchase ~ ., data =
Caravan.train, n.trees = 1000, shrinkage = 0.01,
distribution = "bernoulli")
Warning: variable 50: PVRAAUT has no
variation.
Warning: variable 71: AVRAAUT has no
variation.
summary(boost.caravan)



```
##          var rel.inf
## PERSAUT PERSAUT 15.15534
## MKOOPKLA MKOOPKLA 9.23500
## MOPLHOOG MOPLHOOG 8.67017
## MBERMIDD MBERMIDD 5.39404
```

```
## MGODGE    MGODGE 5.03048
....
## AVRAAUT   AVRAAUT 0.00000
## AAANHANG  AAANHANG 0.00000
## ATTRACTOR ATTRACTOR 0.00000
## AWERKT    AWERKT 0.00000
## ABROM     ABROM 0.00000
## ALEVEN    ALEVEN 0.00000
## APERSONG  APERSONG 0.00000
## AGEZONG   AGEZONG 0.00000
## AWAOREG   AWAOREG 0.00000
## AZEILPL   AZEILPL 0.00000
## APLEZIER  APLEZIER 0.00000
## AFIETS    AFIETS 0.00000
## AINBOED   AINBOED 0.00000
## ABYSTAND  ABYSTAND 0.00000
```

PPERSAUT, MKOOPKLA and MOPLHOOG are three most important variables in that order.

```
c)
boost.prob = predict(boost.caravan, Caravan.test, n.trees = 1000, type = "response")
boost.pred = ifelse(boost.prob > 0.2, 1, 0)
table(Caravan.test$Purchase, boost.pred)
##   boost.pred
##      0      1
## 0 4396 137
## 1  255   34
34/(137 + 34)
## [1] 0.1988
```

About 20% of people predicted to make purchase actually end up making one.

```
lm.caravan = glm(Purchase ~ ., data = Caravan.train, family = binomial)
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
lm.prob = predict(lm.caravan, Caravan.test, type = "response")
## Warning: prediction from a rank-deficient fit may be misleading
lm.pred = ifelse(lm.prob > 0.2, 1, 0)
table(Caravan.test$Purchase, lm.pred)
##   lm.pred
##      0      1
## 0 4183 350
## 1  231   58
58/(350 + 58)
## [1] 0.1422
```

About 14% of people predicted to make purchase using logistic regression actually end up making one. This is lower than boosting.